

## WEEK 3 AI FOR SOFTWARE ENGINEERING ASSIGNMENT

### PART 1

#### QUESTION 1

**Q1. The primary differences between TensorFlow and PyTorch and when to choose one over the other.**

##### **1. Programming Style**

- **PyTorch:** Imperative (eager execution) – you write code and execute it line by line. Feels more "Pythonic" and intuitive.
- **TensorFlow:** Originally static graph-based (define-and-run), though **eager execution** is now supported (like PyTorch).

**Choose PyTorch** if you prefer easier debugging and a more intuitive coding experience.

---

##### **2. Model Deployment**

- **TensorFlow:** Strong deployment ecosystem – supports **TensorFlow Lite**, **TensorFlow Serving**, **TensorFlow.js**, and **TFHub**.
- **PyTorch:** Uses **TorchServe** and **ONNX**, but its ecosystem for deployment is relatively newer.

**Choose TensorFlow** if you want seamless **cross-platform model deployment** (mobile, web, embedded).

---

##### **3. Community and Ecosystem**

- **TensorFlow:** Backed by **Google**; has a mature ecosystem and extensive tooling (e.g., **TensorBoard**, **TFX**).
- **PyTorch:** Backed by **Meta (Facebook)**; growing rapidly in research and production usage.

**Choose PyTorch** if you're in **academic/research settings**, TensorFlow for **production-grade enterprise solutions**.

---

##### **4. Model Serving and Production Pipelines**

- **TensorFlow:** More robust tools for end-to-end production pipelines (TFX, Data Validation, etc.).
- **PyTorch:** Easier prototyping, but production support catching up.

**Choose TensorFlow** for **enterprise-scale MLOps**; PyTorch for fast prototyping and iterative development.

---

## 5. Performance

- Both have high performance and GPU acceleration via CUDA.
- TensorFlow has better **multi-GPU and TPU support** out of the box.

**Choose TensorFlow** if you're leveraging **TPUs or distributed training**.

---

### Summary Decision Guide:

Use Case	Recommendation
Rapid prototyping & research	<b>PyTorch</b>
Large-scale production deployment	<b>TensorFlow</b>
Beginners learning deep learning	<b>PyTorch</b>
Mobile or embedded AI apps	<b>TensorFlow</b>
Cutting-edge AI research papers	<b>PyTorch</b>

## Q2. Two use cases for Jupyter Notebooks in AI development.

### Introduction

Jupyter Notebooks offer an interactive environment that seamlessly integrates code, data visualizations, and narrative text, making them invaluable for AI development. Below are two detailed use cases:

---

### 1. Exploratory Data Analysis (EDA) and Preprocessing

#### Overview:

When building AI models, understanding and preparing your data is crucial. Jupyter Notebooks facilitate exploratory data analysis by allowing you to:

- **Interactive Data Exploration:**  
Load and view datasets using libraries like Pandas and NumPy, and then generate

dynamic visualizations with tools such as Matplotlib, Seaborn, or Plotly. This interactivity helps in identifying patterns, outliers, or trends within the data.

- **Step-by-Step Data Cleaning:**  
Experiment with different data cleaning techniques interactively (e.g., handling missing data, normalization, feature scaling, and feature selection). You can iteratively test and refine your preprocessing pipeline within the same document, which is invaluable for tailoring your dataset to the needs of your AI model.
- **Documentation and Sharing:**  
Combine code with markdown cells to document insights, hypotheses, and methodology. This makes it easier to share your findings with team members or stakeholders and maintain reproducibility of your data preparation process.

#### When It's Useful:

- During the early stages of model development, where understanding the data is key.
  - For collaborative projects, where clear communication of the data cleaning steps and insights is needed.
- 

## 2. Prototyping, Experimenting, and Model Iteration

### Overview:

Jupyter Notebooks serve as an ideal platform for rapid prototyping and iterative experimentation with AI models. They allow you to:

- **Model Development and Experimentation:**  
Write and execute code in small, manageable blocks, enabling you to quickly develop, test, and refine machine learning models. Libraries such as TensorFlow, PyTorch, and Scikit-learn can be integrated directly into your notebook cells.
- **Visualization of Model Outputs:**  
Immediately visualize the predictions, loss curves, and performance metrics as your model trains or when comparing different models. This immediate feedback loop is crucial for adjusting hyperparameters, diagnosing issues like overfitting, and enhancing overall model performance.
- **Integration of Narrative and Code:**  
Annotate experiments with explanations, record observations, and highlight performance improvements using markdown cells. This helps create a clear narrative of the model development process and serves as a reference for future iterations.
- **Reproducibility and Collaboration:**  
Jupyter Notebooks can be version-controlled and shared among colleagues. This collaborative approach ensures that each experiment is documented and can be reproduced or built upon by other team members, fostering an efficient and creative development cycle.

#### When It's Useful:

- In research or academic settings where iterative experimentation and documentation are vital.
  - In industry projects that benefit from rapid prototyping and collaborative debugging sessions.
- 

## Conclusion

These two use cases highlight the flexibility and power of Jupyter Notebooks in AI development, supporting both the foundational tasks of data exploration and the dynamic requirements of model prototyping and iteration.

## Q3. How spaCy enhance NLP tasks compared to basic Python string operations.

### Introduction

spaCy significantly enhances NLP (Natural Language Processing) tasks compared to basic Python string operations by providing **linguistically intelligent**, **efficient**, and **scalable** tools for analyzing and processing human language. Here's how:

---

### 1. Tokenization and Text Structure Understanding

- **Basic Python:** Uses `split ()` or `regex`, which splits text naively by whitespace or punctuation.
- **spaCy:** Performs *intelligent tokenization* — understanding contractions (e.g., "can't" → "ca" + "n't"), punctuation, and special cases using built-in language models.

**Benefit:** More accurate parsing of real-world language structures.

---

### 2. Part-of-Speech (POS) Tagging

- **Basic Python:** Has no built-in POS tagging capability.
- **spaCy:** Automatically labels each word with its grammatical role (e.g., noun, verb, adjective).

**Benefit:** Enables syntax-aware processing — useful in tasks like grammar checking, information extraction, or question answering.

---

### 3. Named Entity Recognition (NER)

- **Basic Python:** Requires manual rule-based approaches or keyword matching.
- **spaCy:** Identifies and classifies named entities (e.g., people, organizations, dates) out of the box.

**Benefit:** Extracts meaningful data (e.g., “Barack Obama” → PERSON) with high accuracy and minimal effort.

---

### 4. Dependency Parsing

- **Basic Python:** Lacks understanding of grammatical relationships.
- **spaCy:** Constructs dependency trees to show how words relate (e.g., subject → verb → object).

**Benefit:** Crucial for complex NLP tasks like semantic role labeling, question generation, or text summarization.

---

### 5. Lemmatization

- **Basic Python:** Requires custom dictionaries or stemming libraries.
- **spaCy:** Reduces words to their base forms (e.g., "running" → "run") using linguistic context.

**Benefit:** Improves search, matching, and classification performance by normalizing text.

---

### 6. Pretrained Models and Language Support

- **Basic Python:** Manual implementation for every task.
- **spaCy:** Comes with pretrained models for multiple languages and tasks, optimized for performance and accuracy.

**Benefit:** Quick access to powerful NLP pipelines without building from scratch.

---

## Summary Comparison

Task	Basic Python	spaCy
Tokenization	str.split() / regex	Context-aware and accurate
POS Tagging	Not available	Built-in with pretrained models
Named Entity Recognition	Manual	Automatic, high-accuracy
Dependency Parsing	Not available	Built-in syntactic parsing
Lemmatization	Manual or NLTK-based	Fast and linguistically correct
Performance	Slower with large text	Highly optimized for speed & scale

### Conclusion:

spaCy elevates text processing from raw string manipulation to true language understanding, enabling more accurate and powerful NLP applications like chatbots, text classifiers, summarizers, and knowledge extractors.

## QUESTION 2

### 1. Target Applications

Aspect	Scikit-learn	TensorFlow
Focus Area	Classical Machine Learning	Deep Learning / Neural Networks
Common Use Cases	Regression, classification, clustering, dimensionality reduction, model evaluation	Image processing, NLP, speech recognition, large-scale neural networks
Model Types	Decision trees, SVMs, random forests, k-NN, logistic regression	CNNs, RNNs, Transformers, GANs, custom DNNs

- Use **Scikit-learn** for traditional ML tasks on structured/tabular data.
- Use **TensorFlow** for complex tasks like computer vision, natural language processing, and large-scale AI.

## 2. Ease of Use for Beginners

Aspect	Scikit-learn	TensorFlow
API Design	Simple, consistent, and intuitive	More complex, especially in low-level APIs
Learning Curve	Gentle – great for beginners	Steeper – especially before TensorFlow 2.0
Workflow Simplicity	Fit/predict pipeline is beginner-friendly	Keras API improves ease, but deep learning remains more complex

- **Scikit-learn is easier** for beginners with basic programming or data science backgrounds.
  - **TensorFlow (via Keras)** can also be accessible but is better for learners with some understanding of neural networks.
- 

## 3. Community Support

Aspect	Scikit-learn	TensorFlow
Community Size	Large and well-established in academia	Very large, especially in industry and research
Documentation	Excellent and beginner-friendly	Extensive, especially with TensorFlow 2.x and Keras
Ecosystem	Integrates well with Pandas, NumPy, etc.	Strong ecosystem: TensorBoard, TF Lite, TF Hub, etc.
Online Resources	Abundant tutorials and examples	Very active forums, GitHub issues, and courses

- **Both have strong communities**, but TensorFlow has broader industrial backing and tools for production deployment.
  - Scikit-learn is **more common in traditional ML academic settings**, while TensorFlow is a **powerhouse for AI in production**.
-

## 🔍 Summary Table:

Feature	Scikit-learn	TensorFlow
Target Applications	Classical ML	Deep Learning & AI
Best for Beginners	Very beginner-friendly	Moderate (easier with Keras)
Community Support	Strong in ML & data science	Massive in AI and industry

## Recommendation:

- Start with **Scikit-learn** if you're new to ML and working with structured data.
- Move to **TensorFlow** when you're ready to build and train deep learning models for complex tasks like image recognition or NLP.



## SCREENSHORTS

### Named Entity Recognition (NER):

Review 1: I love my new Samsung Galaxy S21, the camera is amazing!  
- Samsung Galaxy S21 (ORG)  
-----

Review 2: The Apple AirPods Pro are so comfortable and the sound is top-notch.  
- No recognizable product/brand entity found.  
-----

Review 3: I was disappointed with the Sony headphones, poor noise cancellation.  
- Sony (ORG)  
-----

Review 4: This HP laptop is very fast and reliable for work and school.  
- HP (ORG)  
-----

Review 5: Avoid the fake Ray-Ban sunglasses - cheap material and broke in 2 days!  
- No recognizable product/brand entity found.  
-----

Review 6: The Dell XPS 13 is sleek and fast, I'm very happy with my purchase.  
- No recognizable product/brand entity found.  
-----

Review 7: Logitech mouse stopped working after just a week. Not worth the price.  
- Logitech (ORG)  
-----

Review 8: Great performance from the ASUS ROG laptop - it runs all my games smoothly.  
- ASUS (ORG)  
-----

Review 9: The Beats Studio3 have excellent bass and look stylish.  
- No recognizable product/brand entity found.  
-----

Review 10: I regret buying the Lenovo tablet. It's very slow and laggy.  
- Lenovo (ORG)  
-----

### Rule-Based Sentiment Analysis:

Review 1: I love my new Samsung Galaxy S21, the camera is amazing!  
- Polarity Score: 0.46 → Sentiment: Positive  
-----

Review 2: The Apple AirPods Pro are so comfortable and the sound is top-notch.  
- Polarity Score: 0.60 → Sentiment: Positive  
-----

Review 3: I was disappointed with the Sony headphones, poor noise cancellation.  
- Polarity Score: -0.57 → Sentiment: Negative  
-----

Review 4: This HP laptop is very fast and reliable for work and school.  
- Polarity Score: 0.26 → Sentiment: Positive  
-----

Review 5: Avoid the fake Ray-Ban sunglasses - cheap material and broke in 2 days!  
- Polarity Score: 0.00 → Sentiment: Neutral  
-----

Review 6: The Dell XPS 13 is sleek and fast, I'm very happy with my purchase.  
- Polarity Score: 0.60 → Sentiment: Positive  
-----

Review 7: Logitech mouse stopped working after just a week. Not worth the price.  
- Polarity Score: -0.15 → Sentiment: Negative  
-----

Review 8: Great performance from the ASUS ROG laptop - it runs all my games smoothly.  
- Polarity Score: 0.60 → Sentiment: Positive  
-----

Review 9: The Beats Studio3 have excellent bass and look stylish.  
- Polarity Score: 0.45 → Sentiment: Positive  
-----

Review 10: I regret buying the Lenovo tablet. It's very slow and laggy.  
- Polarity Score: -0.39 → Sentiment: Negative  
-----

Logistic Regression Accuracy: 1.0  
Decision Tree Accuracy: 1.0

### Decision Tree Visualization

